

Leveraging Modern C++ in Quantitative Finance


CppCon
September 2019

Daniel Hanson
Lecturer, Computational Finance & Risk Management
Dept of Applied Mathematics
University of Washington




What this talk is about

- Focus on end-users
- Easy to use, powerful math-related tools in C++ we can *leverage*



The Start of Something Beautiful

- C++11
 - Math features (coming up)
 - Move semantics => more efficient models code
 - Parallel STL algorithms followed in C++17
- Better availability of quality open source math libraries (late 2000's)
 - Eigen
 - Armadillo
 - More recently:
 - xtensor – a la NumPy (~2016)
 - DataFrame – a la R dataframe and Python Pandas (~2018)
- Coming attractions
 - A proper Date class with C++20
 - SG14 linear algebra library proposal



- Mathematical contributions to Boost
 - Some are very intuitive and useful
 - But others could be a lot more user-friendly


MATH IN THE NEW C++ STANDARD

C++11 and After

Case study: Monte Carlo model for pricing a European option

Monte Carlo Option Pricing

- European Option: A tradeable contract that gives the holder the right to buy or sell a share of stock at a predetermined strike price on its expiration date
- A simple graphical example of a call option is shown below (right to purchase at strike price)
 - Assume the vertical axis represents changes from an underlying asset currently valued at \$100/share
 - The red line represents a strike price of \$105
 - The positive payoffs at expiration will be the blue (~\$15), green (~\$10) and red (~\$5) scenarios



Monte Carlo Option Pricing

- In our five-scenario example, to determine the option price:
 - If the risk-free interest rate is, say 1.2%,
 - and the time to expiration is four months (1/3 of a year),
 - the option value would be

$$\frac{e^{-(0.012)(0.333)} \{ (120 - 105) + (115 - 105) + (110 - 105) + 0 + 0 \}}{5} = 5.98$$


Discount
factor

Calculate average to get
option price

Payoffs at
expiration

Monte Carlo Option Pricing

- In reality, the number of scenarios required can be on the order of 10,000 – 100,000
- This can lead to computationally intensive operations
- First, however, we need to generate *a single equity price path* (next slide)



$$S_t = S_{t-1} e^{\left(r - \frac{\sigma^2}{2}\right) \Delta t + \sigma \varepsilon_t \sqrt{\Delta t}}$$

$$S_{t-1} = S_0 = \text{initEquityPrice_}$$

$$r = \text{rfRate_}$$

$$\sigma = \text{volatility_}$$

$$\Delta t = \text{dt_}$$

$\varepsilon_t = N(0, 1)$ random variates to be generated in the code (see next slide)

```
#include "EquityPriceGenerator.h"  
#include <random>  
using std::mt19937_64;  
using std::normal_distribution;
```

Equity Price Generator (implementation)

```
vector<double> EquityPriceGenerator::operator()(int seed) const
{
    vector<double> v;
    mt19937_64 mtEngine(seed);
    normal_distribution<> nd;
    auto newPrice = [this](double previousEquityPrice, double norm)
    {
        double price = 0.0;
        double expArg1 = (rfRate_ - ((volatility_ * volatility_) / 2.0)) * dt_;
        double expArg2 = volatility_ * norm * sqrt(dt_);
        price = previousEquityPrice * exp(expArg1 + expArg2);
        return price;
    };

    v.push_back(initEquityPrice_);

    double equityPrice = initEquityPrice_;
    // i <= numTimeSteps_ since we need a price at the end of the
    // final time step.
    for (int i = 1; i <= numTimeSteps_; ++i)
    {
        equityPrice = newPrice(equityPrice, nd(mtEngine));
        v.push_back(equityPrice);
    }
    return v;
}
```

Where the magic starts

$S_t = S_{t-1} e^{\left(r - \frac{\sigma^2}{2}\right) \Delta t + \sigma \varepsilon_t \sqrt{\Delta t}}$

Lambda Expression implements discretized SDE

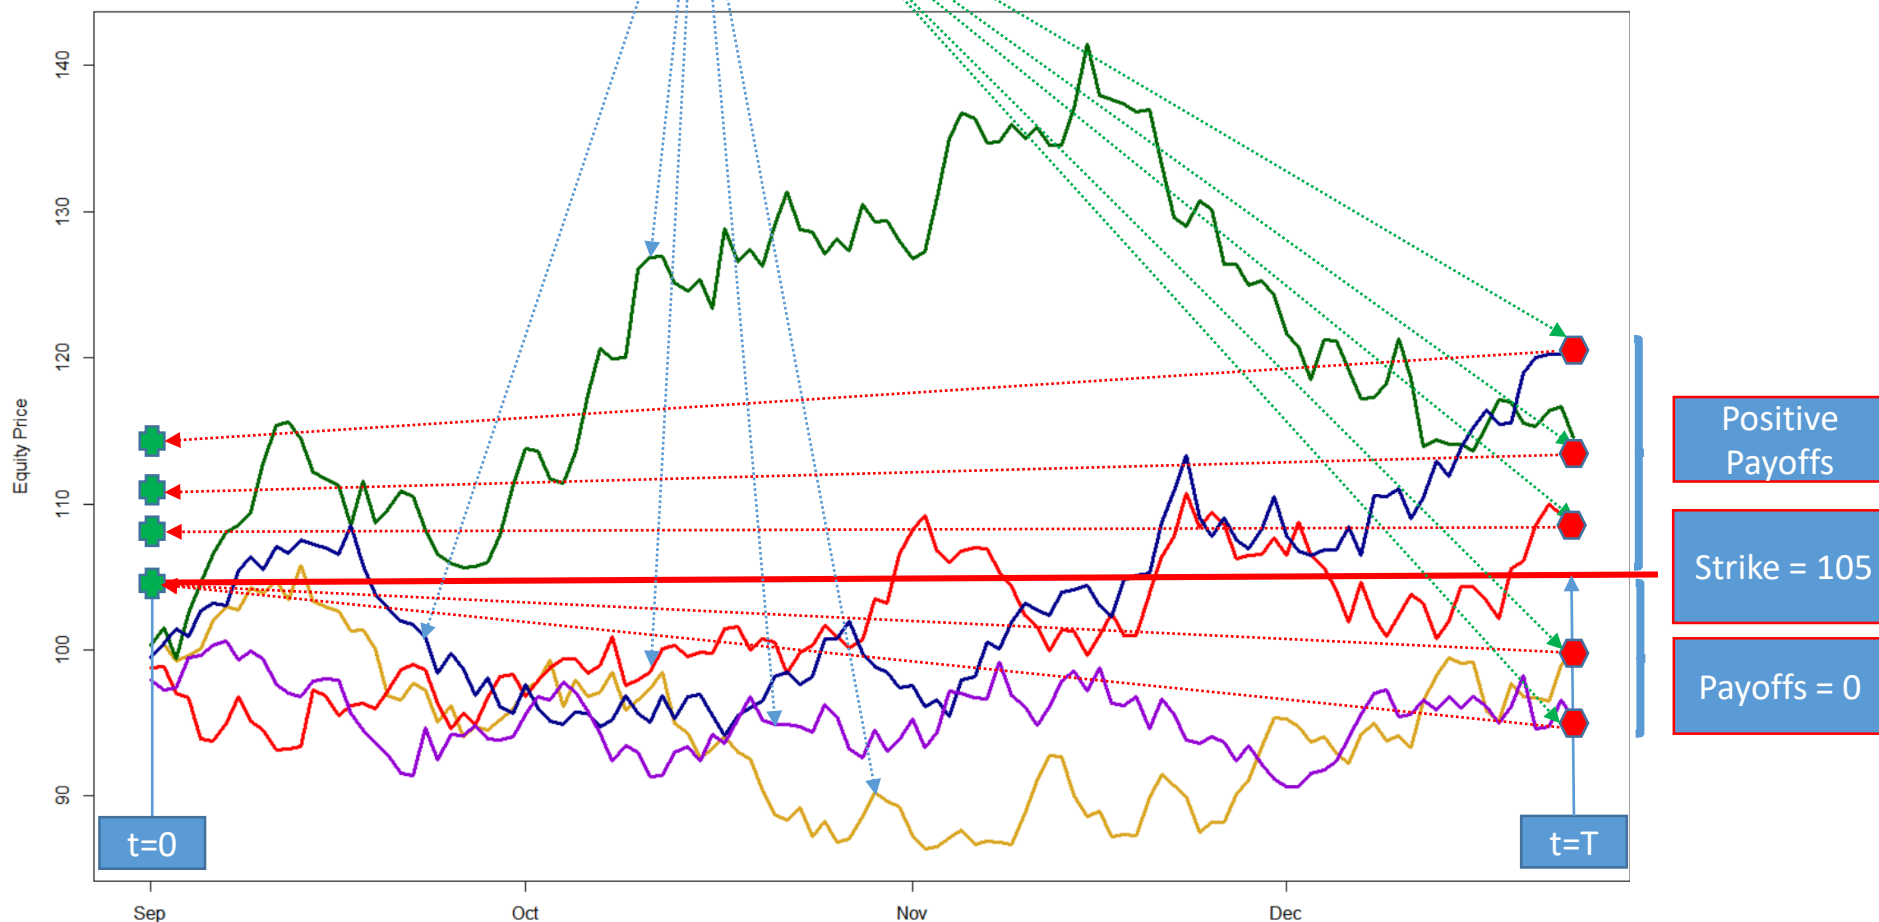
// put initial equity price into
// the 1st position in the vector

S_0

// norm = nd(mtEngine)

Option Pricing with Task-Based Concurrency

- Going back to our five-scenario illustration:
 - Each generated path is a parallel task: `std::future<std::vector<double>>`
 - For each future object, call `get()` to obtain the vector of simulated prices
 - For each vector of prices, call `back()` to get the terminal price
 - Use these terminal prices to compute each discounted payoff \Rightarrow average = option price



Option Calculations (Declaration)

- First look at the function declarations

```
#include "EquityPriceGenerator.h"
```

Use our previous result

```
enum class OptionType
```

```
{  
    CALL,  
    PUT  
};
```

enum class represents the option type: Call or Put

```
class MCEuroOptPricer
```

```
{
```

```
public:
```

Additional option model inputs

```
MCEuroOptPricer(double strike, double spot, double riskFreeRate, double volatility,  
    double timeToExpiry, OptionType optionType, int numTimeSteps, int numScenarios,  
    bool runParallel, int initSeed, double quantity);
```

```
double operator()() const; // Returns the option price
```

```
private:
```

```
// Compare results: non-parallel vs in-parallel with async and futures
```

```
void computePriceNoParallel_();
```

```
void computePriceAsync_();
```

```
// ...
```

```
};
```

computePriceAsync()
will generate each scenario
path in parallel

```

void MCEuroOptPricer::computePriceAsync_(
{
    EquityPriceGenerator epg(spot_, numTimeSteps_, timeToExpiry_, riskFreeRate_, volatility_);
    generateSeeds_();

    using realVector = std::vector<double>;
    std::vector<std::future<realVector> > futures;
    futures.reserve(numScenarios_);

    for (auto& seed : seeds_)
    {
        futures.push_back(std::async(epg, seed));
    }

    realVector discountedPayoffs;
    discountedPayoffs.reserve(numScenarios_);

    for (auto& future : futures)
    {
        double terminalPrice = future.get().back();
        double payoff = payoff_(terminalPrice);
        discountedPayoffs.push_back(discFactor_ * payoff);
    }

    double numScens = static_cast<double>(numScenarios_);
    price_ = quantity_ * (1.0 / numScens) *
        std::accumulate(discountedPayoffs.begin(), discountedPayoffs.end(), 0.0);
}

```

`#include<future>`

Create **vector** of **future** objects

`async` instructs the program to generate a **vector** of simulated stock prices on a separate thread; each threaded operation is encapsulated in a **future** object

The `get()` member function on the **future** object retrieves the asynchronously generated **vector**. We only need the last simulated equity price, so we just use `back()`.


Option Calculations (Summary and Results)

- Hyper-V 20-core virtual machine
- Windows 2019 Server
- Visual Studio 2017 Compiler

Expiry (years)	Time Steps	Scenarios	Time – Serial (sec)	Time - Parallel (Task)	Pct Drop in RT
1	12	10000	0.027	0.027	0.0%
1	120	50000	0.608	0.179	70.6%
5	120	50000	0.604	0.18	70.2%
5	600	50000	2.643	0.261	90.1%
5	600	100000	5.224	0.563	89.2%
10	600	100000	5.17	0.576	88.9%
10	1200	100000	10.201	0.933	90.9%

Option Calculations (Summary and Results)

- This example used the simple case of a European equity option
- Much more complex and computationally intensive option contracts exist
- Remark: There are different engine algorithms and distributions available in `<random>`
 - Mersenne Twister 64 is the most robust engine
 - Total of 17 well-known “textbook” distributions
 - See Nicolai Josuttis: The C++ Standard Library (2E), §17.1.4: Distributions



MATH IN THE BOOST LIBRARIES

- Boost Math Toolkit (2.9.0):
https://www.boost.org/doc/libs/1_70_0/libs/math/doc/html/index.html
 - Statistical Distributions
 - Numerical Integration
- Additional Math-Related Boost Libraries (not in Math Toolkit)
 - Circular Buffers
 - Accumulators
 - MultiArray

Boost Math Toolkit: Statistical Distributions

- Each probability distribution in this library is a class type
- Examples: Construct objects of Student's t and Standard Normal distribution types:

```
#include <boost/math/distributions/students_t.hpp>
```

```
#include <boost/math/distributions/normal.hpp>
```

```
using boost::math::students_t;
```

```
using boost::math::normal;
```

```
// Construct a students_t distribution with 4 degrees of freedom:
```

```
students_t d1(4);
```

```
// Construct a normal distribution with mean 0 and std dev 1:
```

```
normal std_normal(0.0, 1.0);
```

```
// See http://www.boost.org/doc/libs/1\_70\_0/libs/math/doc/html/math\_toolkit/stat\_tut/overview/objects.html
```

Boost Math Toolkit: Statistical Distributions

- We can then apply the *cumulative distribution function*, *probability density function*, and *quantile function* for any distribution.
- Generic operations are non-member functions
- Want to calculate the PDF (Probability Density Function) of a distribution? No problem, just use:
 - `pdf(my_dist, x);` // Returns PDF (density) at point x
// of distribution my_dist.
- Or how about the CDF (Cumulative Distribution Function):
 - `cdf(my_dist, x);` // Returns CDF (integral from -
// infinity to point x)
// of distribution my_dist.
- And quantiles are similar:
 - `quantile(my_dist, p);` // Returns the value of the
// random variable x such that
// `cdf(my_dist, x) == p`.

- 34 probability distributions available in Boost > 17 in Standard Lib
- Will we see a union of all these distributions and functions in the Standard Library eventually?

Boost Math Toolkit: Integration and Differentiation

- Example: Trapezoid Rule to evaluate integral that computes π :

```
auto f = [](double x)
{
    return 4.0 / (1.0 + x * x);
};
```

The function `f` can also be represented as a functor

```
double appPi = trapezoidal(f, 0.0, 1.0);    // Boost function (default)
```

3.14159


-
- All numerical differentiation and integration methods will accept the function input as a *lambda* or a *functor*
 - Could also be applied to root finding algorithms; examples in sample code:
 - Bisection Method
 - Steffenson's Method

Some Additional Math-Related Boost Libraries (outside Math Toolkit)

- Circular Buffers (CircularBuffers.cpp in sample code)
 - STL compliant
 - Similar to `std::deque`, but with fixed capacity
 - Very useful for managing dynamic time series data
 - Old data popped off at capacity as new data pushed on at back
- Accumulators (Accumulators.cpp in sample code)
 - STL compliant
 - Ideal for managing data columns
 - Functions for descriptive statistics (min, max, mean, median, variance, etc) come implemented – very convenient

Some Additional Math-Related Boost Libraries (outside Math Toolkit)

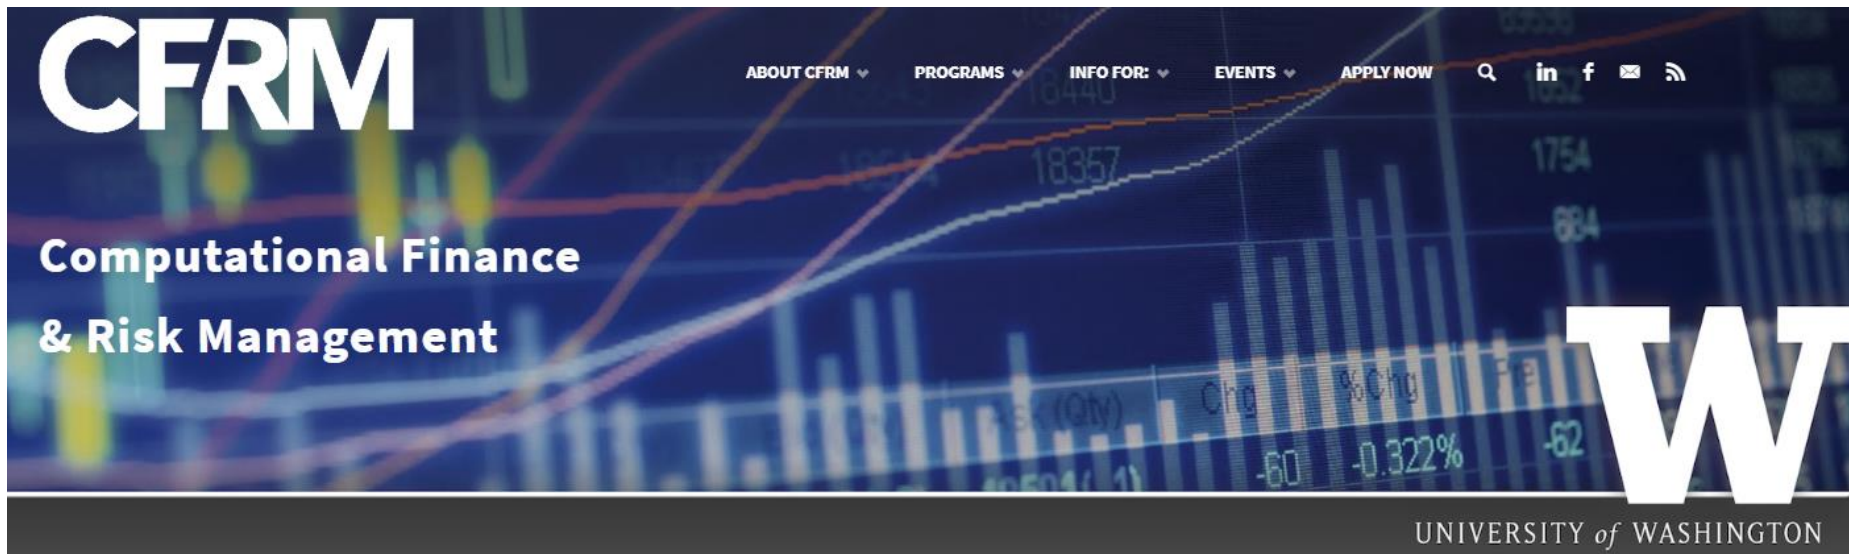
- MultiArray (MultiArray.cpp in sample code)
 - STL compliant
 - Templated multidimensional array
 - Useful for lattice models for pricing options
 - Binomial lattice for option on single asset



Peter James: Option Theory, Figure 7.4 (left), Figure 12.1 (right)

References

- Jon Kalb and Gasper Azman, C++ Today: The Beast is Back, O'Reilly
- Nicolai Josuttis, The C++ Standard Library (2E), Addison Wesley
- Scott Meyers, Effective Modern C++, O'Reilly
- Peter James, Option Theory, Wiley
- Numerical Library in C++ project (2019): Root finding algorithms implemented by Tania Luo, MSc student, Dept of Applied Mathematics, Univ of Washington
- Information about the CFRM program at the University of Washington: <http://cfrm.uw.edu>



- Accompanying sample code available on GitHub:
<https://github.com/QuantDevHacks/CppCon2019Backup>
- Contact:
 - email: [hansondj \(at\) uw.edu](mailto:hansondj@uw.edu)
 - LinkedIn: <https://www.linkedin.com/in/danielhanson/>
- Leverage a beverage



- Thanks for attending!